

4

Solutions

4.1

4.1.1 The value of the signals is as follows:

RegWrite	ALUSrc	ALUoperation	MemWrite	MemRead	MemToReg
true	0	"and"	false	false	0

Mathematically, the MemRead control wire is a "don't care": the instruction will run correctly regardless of the chosen value. Practically, however, MemRead should be set to false to prevent causing a segment fault or cache miss.

4.1.2 Registers, ALUSrc mux, ALU, and the MemToReg mux.

4.1.3 All blocks produce some output. The outputs of DataMemory and the Sign Extender are not used.

4.2 Reg2Loc for LDUR: When executing LDUR, it doesn't matter which value is passed to "Read register 2", because the ALUSrc mux ignores the resulting "Read data 2" output and selects the sign extended immediate value instead.

MemToReg for STUR and CBZ: Neither STUR nor CBZ write a value to the register file. It doesn't matter which value the MemToReg mux passes to the register file because the register file ignores that value.

4.3

4.3.1 $25 + 10 = 35\%$. Only LDUR and STUR use Data memory.

4.3.2 100% Every instruction must be fetched from instruction memory before it can be executed.

4.3.3 $28 + 25 + 10 + 11 + 2 = 76\%$. Only R-type instructions do not use the Sign extender.

4.3.4 The sign extend produces an output during every cycle. If its output is not needed, it is simply ignored.

4.4

4.4.1 Only LDUR is broken. MemToReg is either 1 or "don't care" for all other instructions.

4.4.2 I-type, LUDR, STUR are all broken.

4.4.3 STUR, CBZ, and CBNZ are all broken.

4.5 For context: The encoded instruction is STUR X2, [X3, 0x14]

4.5.1

Sign-Extend	Shift left 2
0x0000000000000014	0x0000000000000050

4.5.2

ALUop	OpCode
00	111 1100 0000 (0x7c0)

4.5.3 The new PC is the old PC + 4. This signal goes from the PC, through the “PC + 4” adder, through the “branch” mux, and back to the PC.

4.5.4 Reg2Loc: Inputs: 1 and 2 and selector 1; Output: 2 (The “upper” input to this mux is bits 20:16 from the instruction. In STUR, bits 20:12 are the offset—in this case, 0x14. Thus, bits 20:16 contain the second hex digit: 0x1.)

ALUsrc: Inputs: Reg[X2] and 0x0000000000000014; Output: 0x0000000000000014

MemToReg: Inputs: Reg[X2] + 0x14 and <undefined>; output: <undefined>/

Branch: Inputs: PC+4 and 0x0000000000000050

4.5.5 ALU inputs: Reg[X3] and 0x0000000000000014

PC + 4 adder inputs: PC and 4

Branch adder inputs: PC and 0x0000000000000050

4.5.6 Read register 1: 3

Read register 2: 2

Write register: 2

Write data: undefined. Either Reg[X3]+0x14, or the undefined output of data memory, depending on the value of the MemToReg control wire.

4.6

4.6.1 No additional logic blocks are needed.

4.6.2 Reg2Loc: X

Uncondbranch: false

Branch: false

MemRead: false (See footnote from solution to problem 4.1.1.)

MemToReg: 0

ALUop: 1X (or simply saying “add” is sufficient for this problem)

MemWrtie: false

ALUsrc: 1

RegWrite: 1

The value of Reg2Loc is a “don’t care” because the ALUsrc mux ignores the resulting “Read data 2” output.

4.7

4.7.1 Figure 2.20 shows that instruction bit 28 contains the correct value for the `Reg2Loc` control line. In general, the value of the `Reg2Loc` wire is 0 for R-type instructions, 1 for D-type and CB-type instructions, and “don’t care” for everything else. `Reg2Loc` is also a “don’t care” for LSL and LSR because, even though they use an R-type format, they don’t use a second register operand.

4.7.2 R-type: $30 + 250 + 25 + 150 + 25 + 200 + 25 + 20 = 725\text{ps}$

4.7.3 LDUR: $30 + 250 + 150 + 200 + 250 + 25 + 20 = 925\text{ ps}$

4.7.4 STUR: $30 + 250 + 150 + 200 + 250 = 880$

4.7.5 CBZ: $30 + 250 + 25 + 150 + 25 + 200 + 5 + 5 + 25 + 20 = 735$

4.7.6 B: $30 + 250 + 50 + 150 + 25 + 20 = 525\text{ps}$

4.7.7 I-type: $30 + 250 + 150 + 200 + 25 + 20 = 675\text{ps}$

4.7.8 925ps

4.8 Using the results from Problem 4.7, we see that the average time per instruction is

$$.52 \cdot 725 + .25 \cdot 925 + .1 \cdot 880 + .11 \cdot 735 + .02 \cdot 525 = 787.6\text{ps}$$

In contrast, a single-cycle CPU with a “normal” clock would require a clock cycle time of 925.

Thus, the speedup would be $925/787.6 = 1.174$

4.9

4.9.1 Without improvement: 925; With improvement: 1225

4.9.2 The running time of a program on the original CPU is $925 \cdot n$. The running time on the improved CPU is $1225 \cdot (0.95) \cdot n = 1163.75$. Thus, the “speedup” is 0.794. (Thus, this “improved” CPU is actually slower than the original).

4.9.3 Because adding a multiply instruction will remove 5% of the instructions, the cycle time can grow to as much as $925/(0.95) = 973.7$. Thus, the time for the ALU can increase by up to 48 (from 200 to 248).

4.10

4.10.1 The additional registers will allow us to remove 12% of the loads and stores, or $(0.12) \cdot (0.25 + 0.1) = 4.2\%$ of all instructions. Thus, the time to run n instructions will decrease from $925 \cdot n$ to $935 \cdot .958 \cdot n = 895.73 \cdot n$. That corresponds to a speedup of $925/895.73 = 1.03$.

4.10.2 The cost of the original CPU is 4507; the cost of the improved CPU is 4707.

PC: 5

I-Mem: 1000

Register file: 200

ALU: 100

D-Mem: 2000

Sign Extend: 100

2 Controls: 1000

2 adders: 30×2

4 muxes: 4×10

2 single gates: 2×1

Thus, for a 3% increase in performance, the cost of the CPU increases by about 4.4%.

4.10.3 From a strictly mathematical standpoint it does not make sense to add more registers because the new CPU costs more per unit of performance. However, that simple calculation does not account for the utility of the performance. For example, in a real-time system, a 3% performance may make the difference between meeting or missing deadlines. In which case, the improvement would be well worth the 4.4% additional cost.

4.11

4.11.1 No new functional blocks are needed.

4.11.2 Only the control unit needs modification.

4.11.3 No new data paths are needed.

4.11.4 No new signals are needed.

4.12

4.12.1 No new functional blocks are needed.

4.12.2 The register file needs to be modified so that it can write to two registers in the same cycle.

4.12.3 There would need to be a data path from read data 1 to write data 2. (There is an ALU op that causes read data 2 to be passed unmodified as output. Thus, in this sense, there is already a path from read data 2 to write data 1.)

4.12.4 There would need to be a second RegWrite control wire.

4.12.5 Many possible solutions.

4.13

4.13.1 We need some additional muxes to drive the data paths discussed in 4.13.3.

4.13.2 No functional blocks need to be modified.

4.13.3 There needs to be a path from the ALU output to data memory's write data port. There also needs to be a path from read data 2 directly to Data memory's Address input.

4.13.4 These new data paths will need to be driven by muxes. These muxes will require control wires for the selector.

4.13.5 Many possible solutions.

4.14 Although many instructions use the sign-extend block, it is on the critical path of the B instruction only. All other instructions that use sign extend also use the register file, which is slower.

4.15

4.15.1 The new clock cycle time would be 725. ALU and Data Memory will now run in parallel, so we have effectively removed the faster of the two (the ALU with time 200) from the critical path.

4.15.2 Slower. The original CPU takes $925 \cdot n$ picoseconds to run n instructions. The same program will have approximately $1.35 \cdot n$ instructions when compiled for the new machine. Thus, the time on the new machine will be $725 \cdot 1.35n = 978.75 \cdot n$. This represents a "speedup" of .95.

4.15.3 The number of loads and stores is the primary factor. How the loads and stores are used can also have an effect. For example, a program whose loads and stores tend to be to only a few different address may also run faster on the new machine.

4.15.4 This answer is a matter of opinion.

4.16

4.16.1 Pipelined: 350; non-pipelined: 1250

4.16.2 Pipelined: 1250; non-pipelined: 1250

4.16.3 Split the ID stage. This reduces the clock-cycle time to 300ps.

4.16.4 35%.

4.16.5 65%

4.17 $n + k - 1$. Let's look at when each instruction is in the WB stage. In a k -stage pipeline, the 1st instruction doesn't enter the WB stage until cycle k . From that point on, at most one of the remaining $n - 1$ instructions is in the WB stage during every cycle.

This gives us a minimum of $k + (n - 1) = n + k - 1$ cycles.

4.18 $X3 = 33$ and $X4 = 36$

4.19 $X5 = 54$ (The code will run correctly because the result of the first instruction is written back to the register file at the beginning of the 5th cycle, whereas the final instruction reads the updated value of $X1$ during the second half of this cycle.)

4.20

```

ADDI X1, X2, 5
NOP
NOP
ADD X3, X1, X2
ADDI X4, X1, 15
NOP
ADD X5, X3, X2

```

4.21

4.21.1 Pipeline without forwarding requires $1.4 \cdot n \cdot 250\text{ps}$. Pipeline with forwarding requires $1.05 \cdot n \cdot 300\text{ps}$. The speedup is therefore $(1.4 \cdot 250) / (1.05 \cdot 300) = 1.11$.

4.21.2 Our goal is for the pipeline with forwarding to be faster than the pipeline without forwarding. Let y be the number of stalls remaining as a percentage of "code" instructions. Our goal is for $300 \cdot (1+y) \cdot n < 250 \cdot 1.4 \cdot n$. Thus, y must be less than 16.7%.

4.21.3 This time, our goal is for $300(1 + y) \cdot n < 250(1 + x) \cdot n$. This happens when $y < (250x - 50)/300$.

4.21.4 It cannot. In the best case, where forwarding eliminates the need for every NOP, the program will take time $300 \cdot n$ to run on the pipeline with forwarding. This is slower than the $250 \cdot 1.075 \cdot n$ required on the pipeline with no forwarding.

4.21.5 Speedup is not possible when the solution to 4.21.3 is less than 0. Solving $0 < (250x - 50)/300$ for x gives that x must be at least 0.2.

4.22**4.22.1** Stalls are marked with ***

```

STUR X16, [X6, #12]  IF ID EX ME WB
LDUR X16, [X6, #8]   IF ID EX ME WB
SUB  X7, X5, X4       IF ID EX ME WB
CBZ  X7, Loop         ** ** IF ID EX ME WB
ADD  X5, X1, X4       IF ID EX ME WB
SUBS X5, X15, X4      IF ID EX ME WB

```

4.22.2 Reordering code won't help. Every instruction must be fetched; thus, every data access causes a stall. Reordering code will just change the pair of instructions that are in conflict.

4.22.3 You can't solve this structural hazard with NOPs, because even the NOPs must be fetched from instruction memory.

4.22.4 35%. Every data access will cause a stall.

4.23

4.23.1 The clock period won't change because we aren't making any changes to the slowest stage.

4.23.2 Moving the MEM stage in parallel with the EX stage will eliminate the need for a cycle between an LDUR and operations that use the result of the LDUR. This can potentially reduce the number of stalls in a program.

4.23.3 Removing the offset from LDUR and STUR may increase the total number of instructions because some LDUR and STUR instructions will need to be replaced with a ADD/LDUR or ADD/STUR pair.

4.24 The second one. A careful examination of Figure 4.59 shows that the need for a stall is detected during the ID stage. It is this stage that prevents the fetch of a new instruction, effectively causing the add to repeat its ID stage.

4.25

4.25.1 ... indicates a stall. ! indicates a stage that does not do useful work.

```

LDUR X10, [X1, #0]  IF ID EX ME | WB
LDUR X11, [X2, #8]   IF ID EX | ME WB
ADD  X12, X10, X11   IF ID | .. EX ME! WB
SUBI X1, X1, #16      IF | .. ID EX ME! WB
CBNZ X12, LOOP       | .. IF ID EX ME! WB!

```

```

LDUR X10, [X1, #0]           IF ID EX ME WB
LDUR X11, [X2, #8]           IF ID EX ME WB
ADD  X12, X10, X11           IF ID .. EX | ME! WB
SUBI X1, X1, #16              IF .. ID | EX ME! WB
CBNZ X12, LOOP                IF | ID EX ME! WB!
Completely busy               | N N N N N N N N |

```


4.25.2 In a particular clock cycle, a pipeline stage is not doing useful work if it is stalled or if the instruction going through that stage is not doing any useful work there. As the diagram above shows, there are not any cycles during which every pipeline stage is doing useful work.

4.26

4.26.1 EX to 1st only:

```
ADD X1, X2, X3
ADD X4, X1, X5
ADD X10, X11, X12
```

MEM to 1st only:

```
LDUR X1, [X2, 0]
ADD X5, X1, X3
ADD X10, X11, X12
```

EX to 2nd only:

```
ADD X1, X2, X3
ADD X4, X7, X5
ADD X6, X1, X2
```

MEM to 2nd only:

```
LDUR X1, [X2, #0]
ADD X10, X11, X12
ADD X5, X1, X3
```

EX to 1st and EX to 2nd:

```
ADD X1, X2, X3
ADD X4, X1, X5
ADD X6, X1, X2
```

4.26.2 EX to 1st only: 2 stalls

```
ADD X1, X2, X3
stall
stall
ADD X4, X1, X5
ADD X10, X11, X12
```

MEM to 1st only: 2 stalls

```
LDUR X1, [X2, #0]
stall
stall
ADD X5, X1, X3
ADD X10, X11, X12
```

EX to 2nd only: 1 stall

ADD X1, X2, X3

ADD X4, X7, X5

stall

ADD X6, X1, X2

MEM to 2nd only: 1 stall

LDUR X1, [X2, #0]

ADD X10, X11, X12

stall

ADD X5, X1, X3

EX to 1st and EX to 2nd: 2 stalls

ADD X1, X2, X3

stall

stall

ADD X4, X1, X5

ADD X6, X1, X2

4.26.3 Consider this code:

```
LDUR X1, [X5, #0]    # MEM to 2nd --- one stall
ADD  X2, X6, X7      # EX to 1st  --- two stalls
ADD  X3, X1, X2
ADD  X10, X11, X12
```

If we analyze each instruction separately, we would calculate that we need to add 3 stalls (one for a “MEM to 2nd” and two for an “EX to 1st only”). However, as we can see below, we need only two stalls:

```
LDUR X1, [X5, #0]
ADD  X2, X6, X7
stall
stall
ADD  X3, X1, X2
ADD  X10, X11, X12
```

4.26.4 Taking a weighted average of the answers from 4.26.2 gives $0.05 \cdot 2 + 0.2 \cdot 2 + 0.05 \cdot 1 + 0.1 \cdot 1 + 0.1 \cdot 2 = 0.85$ stalls per instruction (on average) for a CPI of 1.85. This means that $0.85/1.85$ cycles, or 46%, are stalls.

4.26.5 The only dependency that cannot be handled by forwarding is from the MEM stage to the next instruction. Thus, 20% of instructions will generate one stall for a CPI of 1.2. This means that 0.2 out of 1.2 cycles, or 17%, are stalls.

4.26.6 If we forward from the EX/MEM register only, we have the following stalls/NOPs

EX to 1st:	0
MEM to 1st:	2
EX to 2nd:	1
MEM to 2nd:	1
EX to 1st and 2nd:	1

This represents an average of $0.05 \cdot 0 + 0.2 \cdot 2 + 0.05 \cdot 1 + 0.10 \cdot 1 + 0.10 \cdot 1 = 0.65$ stalls/instruction. Thus, the CPI is 1.65

If we forward from MEM/WB only, we have the following stalls/NOPs

EX to 1st:	1
MEM to 1st:	1
EX to 2nd:	0
MEM to 2nd:	0
EX to 1st and 2nd:	1

This represents an average of $0.05 \cdot 1 + 0.2 \cdot 1 + 0.1 \cdot 1 = 0.35$ stalls/instruction. Thus, the CPI is 1.35.

4.26.7

	No forwarding	EX/MEM	MEM/WB	Full Forwarding
CPI	1.85	1.65	1.35	1.2
Period	120	120	1.20	130
Time	222n	198n	162n	156n
Speedup	–	1.12	1.37	1.42

4.26.8 CPI for full forwarding is 1.2

CPI for “time travel” forwarding is 1.0

clock period for full forwarding is 130

clock period for “time travel” forwarding is 230

Speedup = $(1.2 \cdot 130) / (1 \cdot 230) = 0.68$ (That means that “time travel” forwarding actually slows the CPU.)

4.26.9 When considering the “EX/MEM” forwarding in 4.26.6, the “EX to 1st” generates no stalls, but “EX to 1st and EX to 2nd” generates one stall. However, “MEM to 1st” and “MEM to 1st and MEM to 2nd” will always generate the same number of stalls. (All “MEM to 1st” dependencies cause a stall, regardless of the type of forwarding. This stall causes the 2nd instruction’s ID phase to overlap with the base instruction’s WB phase, in which case no forwarding is needed.)

4.27

4.27.1 ADD X5, X2, X1
 NOP
 NOP
 LDUR X3, [X5, #4]
 LDUR X2, [X2, #0]
 NOP
 ORR X3, X5, X3
 NOP
 NOP
 STUR X3, [X5, #0]

4.27.2 It is not possible to reduce the number of NOPs.

4.27.3 The code executes correctly. We need hazard detection only to insert a stall when the instruction following a load uses the result of the load. That does not happen in this case.

4.27.4

Cycle	1	2	3	4	5	6	7	8	
ADD	IF	ID	EX	ME	WB				
LDUR		IF	ID	EX	ME	WB			
LDUR			IF	ID	EX	ME	WB		
ORR				IF	ID	EX	ME	WB	
STUR					IF	ID	EX	ME	WB

Because there are no stalls in this code, PCWrite and IF/IDWrite are always 1 and the mux before ID/EX is always set to pass the control values through.

- (1) ForwardA = X; ForwardB = X (no instruction in EX stage yet)
- (2) ForwardA = X; ForwardB = X (no instruction in EX stage yet)
- (3) ForwardA = 0; ForwardB = 0 (no forwarding; values taken from registers)
- (4) ForwardA = 2; ForwardB = 0 (base register taken from result of previous instruction)
- (5) ForwardA = 1; ForwardB = 1 (base register taken from result of two instructions previous)
- (6) ForwardA = 0; ForwardB = 2 (Rn = X5 taken from register; Rm = X3 taken from result of 1st LDUR—two instructions ago)
- (7) ForwardA = 0; ForwardB = 2 (base register taken from register file. Data to be written taken from previous instruction)

4.27.5 The hazard detection unit additionally needs the values of R_d that comes out of the MEM/WB register. The instruction that is currently in the ID stage needs to be stalled if it depends on a value produced by (or forwarded from) the instruction in the EX or the instruction in the MEM stage. So we need to check the destination register of these two instructions if the instruction is an I-type, R-type, or load. (The hazard detection unit already has access to the op code.) The Hazard unit already has the value of R_d from the EX/MEM register as inputs, so we need only add the value from the MEM/WB register.

No additional outputs are needed. We can stall the pipeline using the three output signals that we already have.

The value of R_d from EX/MEM is needed to detect the data hazard between the ADD and the following LDUR. The value of R_d from MEM/WB is needed to detect the data hazard between the first LDUR instruction and the ORR instruction.

4.27.6

Cycle	1	2	3	4	5	6
ADD	IF	ID	EX	ME	WB	
LDUR		IF	ID	–	–	EX
LDUR			IF	–	–	ID

- (1) $PCWrite = 1$; $IF/IDWrite = 1$; control mux = 0
- (2) $PCWrite = 1$; $IF/IDWrite = 1$; control mux = 0
- (3) $PCWrite = 1$; $IF/IDWrite = 1$; control mux = 0
- (4) $PCWrite = 0$; $IF/IDWrite = 0$; control mux = 1
- (5) $PCWrite = 0$; $IF/IDWrite = 0$; control mux = 1

4.28

4.28.1 The CPI increases from 1 to 1.4125.

An incorrectly predicted branch will cause three instructions to be flushed: the instructions currently in the IF, ID, and EX stages. (At this point, the branch

instruction reaches the MEM stage and updates the PC with the correct next instruction.) In other words, 55% of the branches will result in the flushing of three instructions, giving us a CPI of $1 + (1 - 0.45)(0.25)3 = 1.4125$. (Just to be clear: the always-taken predictor is correct 45% of the time, which means, of course, that it is incorrect $1 - 0.45 = 55\%$ of the time.)

4.28.2 The CPI increases from 1 to 1.3375. $(1 + (.25)(1 - .55) = 1.1125)$

4.28.3 The CPI increases from 1 to 1.1125. $(1 + (.25)(1 - .85) = 1.0375)$

4.28.4 The speedup is approximately 1.019.

Changing half of the branch instructions to an ALU instruction reduces the percentage of instructions that are branches from 25% to 12.5%. Because predicted and mispredicted branches are replaced equally, the misprediction rate remains 15%. Thus, the new CPU is $1 + (.125)(1 - .85) = 1.01875$. This represents a speedup of $1.0375 / 1.01875 = 1.0184$

4.28.5 The “speedup” is .91.

There are two ways to look at this problem. One way is to look at the two ADD instruction as a branch with an “extra” cycle. Thus, half of the branches have 1 extra cycle; 15% of the other half have 1 extra cycles (the pipeline flush); and the remaining branches (those correctly predicted) have no extra cycles. This gives us a CPI of $1 + (.5)(.25)*1 + (.5)(.25)(.15)*1 = 1.14375$ and a speedup of $1.0375 / 1.14375 = .91$.

We can also treat the ADD instructions as separate instructions. The modified program now has $1.125n$ instructions (half of 25% produce an extra instruction). $.125n$ of these $1.125n$ instruction (or 11.1%) are branches. The CPI for this new program is $1 + (.111)(.15)*1 = 1.01665$. When we factor in the 12.5% increase in instructions, we get a speedup of $1.0375 / (1.125 * 1.01665) = .91$.

4.28.6 The predictor is 25% accurate on the remaining branches. We know that 80% of branches are always predicted correctly and the overall accuracy is 0.85. Thus, $0.8*1 + 0.2*x = 0.85$. Solving for x shows that $x = 0.25$.

4.29**4.29.1**

Always Taken	Always not-taken
3/5 = 60%	2/5 = 40%

4.29.2

Outcomes	Predictor value at time of prediction	Correct or Incorrect	Accuracy
T, NT, T, T	0,1,0,1	I,C,I,I	25%

4.29.3 The first few recurrences of this pattern do not have the same accuracy as the later ones because the predictor is still warming up. To determine the accuracy in the “steady state”, we must work through the branch predictions until the predictor values start repeating (i.e. until the predictor has the same value at the start of the current and the next recurrence of the pattern).

Outcomes	Predictor value at time of prediction	Correct or Incorrect (in steady state)	Accuracy
T, NT, T, T, NT	1st occurrence: 0,1,0,1,2 2nd occurrence: 1,2,1,2,3 3rd occurrence: 2,3,2,3,3 4th occurrence: 2,3,2,3,3	C,I,C,C,I	60%

4.29.4 The predictor should be an N-bit shift register, where N is the number of branch outcomes in the target pattern. The shift register should be initialized with the pattern itself (0 for NT, 1 for T), and the prediction is always the value in the leftmost bit of the shift register. The register should be shifted after each predicted branch.

4.29.5 Since the predictor’s output is always the opposite of the actual outcome of the branch instruction, the accuracy is zero.

4.29.6 The predictor is the same as in part d, except that it should compare its prediction to the actual outcome and invert (logical NOT) all the bits in the shift register if the prediction is incorrect. This predictor still always perfectly predicts the given pattern. For the opposite pattern, the first prediction will be incorrect, so the predictor’s state is inverted and after that the predictions are always correct. Overall, there is no warm-up period for the given pattern, and the warm-up period for the opposite pattern is only one branch.

4.30**4.30.1**

Instruction 1	Instruction 2
Invalid target address (EX)	Invalid data address (MEM)

- 4.30.2** The Mux that selects the next PC must have inputs added to it. Each input is a constant address of an exception handler. The exception detectors must be added to the appropriate pipeline stage and the outputs of these detectors must be used to control the pre-PC Mux, and also to convert to NOPs instructions that are already in the pipeline behind the exception-triggering instruction.
- 4.30.3** Instructions are fetched normally until the exception is detected. When the exception is detected, all instructions that are in the pipeline after the first instruction must be converted to NOPs. As a result, the second instruction never completes and does not affect pipeline state. In the cycle that immediately follows the cycle in which the exception is detected, the processor will fetch the first instruction of the exception handler.
- 4.30.4** This approach requires us to fetch the address of the handler from memory. We must add the code of the exception to the address of the exception vector table, read the handler's address from memory, and jump to that address. One way of doing this is to handle it like a special instruction that puts the address in EX, loads the handler's address in MEM, and sets the PC in WB.
- 4.30.5** We need a special instruction that allows us to move a value from the (exception) Cause register to a general-purpose register. We must first save the general-purpose register (so we can restore it later), load the Cause register into it, add the address of the vector table to it, use the result as an address for a load that gets the address of the right exception handler from memory, and finally jump to that handler.

4.31.1

[illegible]

```

B.NE TOP
<not filled>
LSL X10, X5, #3
<not filled>
ADD X11, X1, X10
<not filled>
LDUR X12, [X11, #0]
<not filled>
LDUR X13, [X11, #8]
<not filled>
SUB X14, X12, X13
<not filled>
ADD X15, X2, X10
<not filled>
STUR X14, [X15, #0]
ADDI X5, X5, #2
CMP X5, X6
<not filled>
B.NE TOP

```

4.31.2 The original code requires 12 cycles/loop on a 1-issue machine (stalls shown below) and 11 cycles/loop on the 2-issue machine. This gives a speedup of $12/11 = 1.09$. (That's a terrible improvement considering we nearly doubled the amount of hardware.) We know that the code takes 11 cycles/iteration on the 2-issue machine because the first instruction in loop 1 (The LSL) is fetched in cycle 5 and the first instruction in iteration 2 cycle 2 is fetched in cycle 16.

```

MOV    X5, XZR
B      ENT
TOP:
  LSL   X10, X5, #3
  ADD   X11, X1, X10
  LDUR  X12, [X11, #0]
  LDUR  X11, [X11, #8]
  <stall>
  SUB   X14, X12, X13
  ADD   X15, X2, X10
  STUR  X14, [X15, #0]
  ADDI  X5, X5, #2
ENT:   CMP  X5, X6
      <stall>
      B.NE TOP

```

4.31.3 Here is one possibility:

```

CBZ    X6, XZR, DONE
MOV    X5, XZR
B      ENT
TOP:
  LSL   X10, X5, #3
  ADD   X11, X1, X10
  LDUR  X12, [X11, #0]
  LDUR  X13, [X11, #8]
  ADDI  X5, X5, #2
  SUB   X14, X12, X13
  ADD   X15, X2, X10
  CMP   X5, X6
  STUR  X14, [X15, #0]
  B.NE  TOP
DONE:

```

If we switch to a “pointer-based” approach, we can save one cycle/loop.

The code below does this:

```
for (i = 0; i != j; i += 2) {
    *b = *a - *(a+1);
    b+=2;
    a+=2;
}
```

```
CBZ X6, XZR, DONE
MOV X5, XZR
B ENT
TOP:
LDUR X12, [X1, #0]
LDUR X13, [X1, #8]
ADDI X5, X5, #2
SUB X14, X12, X13
STUR X14, [X2, #0]
CMP X5, X6
ADDI X1, X1, #8
ADDI X2, X2, #8
B.NE TOP
DONE:
```

4.31.4 Here is one possibility:

```
CBZ X6, DONE
MOV X5, XZR
TOP:
LSL X10, X5, #3
ADD X11, X1, X10
LDUR X12, [X11, #0]
ADD X15, X2, X10
LDUR X13, [X11, #8]
ADDI X5, X5, #2
CMP X5, X6
SUB X14, X12, X13
STUR X14, [X15, #0]
B.NE TOP
DONE:
```

If we switch to a “pointer-based” approach, we can save one cycle/loop.

The code below does this:

```
for (i = 0; i != j; i += 2) {
    *b = *a - *(a+1);
    b+=2;
    a+=2;
}

    CBZ    X6, DONE
    MOV    X5, XZR
TOP:
    LDUR   X12, [X1, #0]
    ADDI   X5, X5, #2
    LDUR   X13, [X1, #8]
    ADDI   X1, X1, #16
    CMP    X5, X6
    SUB    X14, X12, X13
    STUR   X14, [X2, #0]
    ADDI   X2, X2, #16
    B.NE  TOP
DONE:
```

4.31.5

```

1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21
CBZ X6, DONE          IF ID EX ME WB
<not filled>

MOV X5, XZR           IF ID EX ME WB
<not filled>

LSL X10, X5, #3       IF ID EX ME WB
<not filled>

ADD X11, X1, X10      IF ID EX ME WB
<not filled>

LDUR X12, [X11, #0]   IF ID EX ME WB
ADDI X15, X2, X10     IF ID EX ME WB
LDUR X13, [X11, #8]   IF ID EX ME WB
ADDI X5, X5, #2       IF ID EX ME WB
CMP X5, X6            IF ID EX ME WB
<not filled>

SUB X14, X12, X13     IF ID EX ME WB
<not filled>

STUR X14, [X15, #0]   IF ID EX ME WB
B.NE TOP              IF ID EX ME WB

LSL X10, X5, #3       IF ID EX ME WB
<not filled>

ADD X11, X1, X10      IF ID EX ME WB
<not filled>

LDUR X12, [X11, #0]   IF ID EX ME WB
ADDI X15, X2, X10     IF ID EX ME WB
LDUR X13, [X11, #8]   IF ID EX ME WB
ADDI X5, X5, #2       IF ID EX ME WB
CMP X5, X6            IF ID EX ME WB
<not filled>

SUB X14, X12, X13     IF ID EX ME WB
<not filled>

STUR X14, [X15, #0]   IF ID EX ME WB
B.NE TOP              IF ID EX ME WB

```

4.31.6 The code from 4.31.3 requires 10 cycles per iteration. The code from 4.31.4 requires 7 cycles per iteration. Thus, the speedup is 1.43.

4.31.7 Here is one possibility:

```
CBZ    X6, XZR, DONE
MOV    X5, XZR
```

TOP:

```
LSL    X10, X5, #3
ADD    X11, X1, X10
LDUR   X12, [X11, #0]
LDUR   X13, [X11, #8]
ADD    X15, X2, X10
SUB    X14, X12, X13
STUR   X14, [X15, #0]
LDUR   X12, [X11, #16]
LDUR   X13, [X11, #24]
ADDI   X5, X5, #4
SUB    X14, X12, X13
CMP    X5, X6
STUR   X14, [X15, #16]
B.NE   TOP
```

DONE:

4.31.8 Here is one possibility:

```
CBZ    X6, XZR, DONE
MOV    X5, XZR
```

```
LSL    X10, X5, #3
```

TOP:

```
ADD    X11, X1, X10
LDUR   X12, [X11, #0]
ADD    X15, X2, X10
LDUR   X13, [X11, #8]
ADDI   X5, X5, #4
SUB    X14, X12, X13
LDUR   X12, [X11, #16]
STUR   X14, [X15, #0]
CMP    X5, X6
LDUR   X13, [X11, #24]
LSL    X10, X5, #3
SUB    X14, X12, X13
STUR   X14, [X15, #16]
B.NE   TOP
```

DONE:

4.31.9 The code from 4.30.7 requires 14 cycles per unrolled iteration. This is equivalent to 7 cycles per original iteration. The code from 4.30.4 requires 8 cycles per unrolled iteration. This is equivalent to 4 cycles per original iteration. Thus, the speedup is 1.75.

4.31.10 Here is one possibility:

```

CBZ   X6, XZR, DONE
MOV   X5, XZR

LSL   X10, X5, #3
ADD   X11, X1, X10

TOP:
LDUR  X12, [X11, #0]
ADD   X15, X2, X10

LDUR  X13, [X11, #8]
ADDI  X5, X5, #4

SUB   X14, X12, X13
LDUR  X12, [X11, #16]

STUR  X14, [X15, #0]
CMP   X5, X6

LDUR  X13, [X11, #24]
LSL   X10, X5, #3

SUB   X14, X12, X13
ADD   X11, X1, X10

STUR  X14, [X15, #16]
B.NE TOP
DONE:

```

This change saves one cycle per unrolled iteration, increasing the speedup to 2.0.

4.32

4.32.1 The energy for the two designs is the same: I-Mem is read, two registers are read, and a register is written. We have: $140\text{pJ} + 2 \cdot 70\text{pJ} + 60\text{pJ} = 340\text{pJ}$.

4.32.2 The instruction memory is read for all instructions. Every instruction also results in two register reads (even if only one of those values is actually used). A load instruction results in a memory read and a register write; a store instruction results in a memory write; all other instructions result in at most a single register write. Because the sum of memory read and register write energy is larger than memory write energy, the worst-case instruction is a load instruction. For the energy spent by a load, we have: $140\text{pJ} + 2 \cdot 70\text{pJ} + 60\text{pJ} + 140\text{pJ} = 480\text{pJ}$.

- 4.32.3** Instruction memory must be read for every instruction. However, we can avoid reading registers whose values are not going to be used. To do this, we must add RegRead1 and RegRead2 control inputs to the Registers unit to enable or disable each register read. We must generate these control signals quickly to avoid lengthening the clock cycle time. With these new control signals, a load instruction results in only one register read (we still must read the register used to generate the address), so our change saves 70pJ (one register read) per load. This is a savings of $70/480 = 14.6\%$.
- 4.32.4** B will benefit, because it need not read any registers at all. CBZ, CBN, and I-type instructions will also benefit because they need only read one register.
- 4.32.5** Before the change, the Control unit decodes the instruction while register reads are happening. After the change, the latencies of Control and Register Read cannot be overlapped. This increases the latency of the ID stage and could affect the processor's clock cycle time if the ID stage becomes the longest-latency stage. However, the sum of the latencies for the register read (90ps) and control unit (150ps) are less than the current 250ps cycle time.
- 4.32.6** If memory is read in every cycle, the value is either needed (for a load instruction), or it does not get past the WB Mux (for a non-load instruction that writes to a register), or it does not get written to any register (all other instructions, including branches and stalls). This change does not affect clock cycle time because the clock cycle time must already allow enough time for memory to be read in the MEM stage. It can affect overall performance if the unused memory reads cause cache misses.

The change also affects energy: A memory read occurs in every cycle instead of only in cycles when a load instruction is in the MEM stage. This increases the energy consumption by 140pJ during 75% of the 250ps clock cycles. This corresponds to a consumption of approximately 0.46 Watts (not counting any energy consumed as a result of cache misses).

4.33

- 4.33.1** To test for a stuck-at-0 fault on a wire, we need an instruction that puts that wire to a value of 1 and has a different result if the value on the wire is stuck at zero.

If the least significant bit of the write register line is stuck at zero, an instruction that writes to an odd-numbered register will end up writing to the even-numbered register. To test for this (1) place a value of 10 in X1, 35 in X2, and 45 in X3, then (2) execute ADD X3,X1,X1. The value of X3 is supposed to be 20. If bit 0 of the Write Register input to the registers unit is stuck at zero, the value is written to X2 instead, which means that X2 will be 40 and X3 will remain at 45.

4.33.2 The test for stuck-at-zero requires an instruction that sets the signal to 1; and the test for stuck-at-1 requires an instruction that sets the signal to 0. Because the signal cannot be both 0 and 1 in the same cycle, we cannot test the same signal simultaneously for stuck-at-0 and stuck-at-1 using only one instruction.

The test for stuck-at-1 is analogous to the stuck-at-0 test: (1) Place a value of 10 in X1, 35 in X2, and 45 in X3, then (2) execute `ADD X2,X1,X1`. The value of X2 is supposed to be 20. If bit 0 of the Write Register input to the registers unit is stuck at 1, the value is written to X3 instead, which means that X3 will be 40 and X2 will remain at 35.

4.33.3 The CPU is still usable. The “simplest” solution is to re-write the compiler so that it uses odd-numbered registers only (not that writing compilers is especially simple). We could also write a “translator” that would convert machine code; however, this would be more complicated because the translator would need to detect when two “colliding” registers are used simultaneously and either (1) place one of the values in an unused register, or (2) push that value onto the stack.

4.33.4 To test for this fault, we need an instruction for which MemRead is set to 1, so it has to be LDUR. The instruction also needs to have branch set to 0, which is the case for LDUR. Finally, the instruction needs to have a different result MemRead is incorrectly set to 0. For a load, setting MemRead to 0 results in not reading memory. When this happens, the value placed in the register is “random” (whatever happened to be at the output of the memory unit). Unfortunately, this “random” value can be the same as the one already in the register, so this test is not conclusive.

4.33.5 Only R-type instructions set RegRt to 1. Most R-type instructions would fail to detect this error because reads are non-destructive—the erroneous read would simply be ignored. However, suppose we issued this instruction: `ADD X1, XZR, XZR`. In this case, if MemRead were incorrectly set to 1, the data memory would attempt to read the value in memory location 0. In many operating systems, address 0 can only be accessed by a process running in protected/kernel mode. If this is the case, then this instruction would cause a segmentation fault in the presence of this error.

4.33.6 To test for this fault, we need an instruction that sets the Branch wire to 1, so it has to be the Branch instruction. However, for the Branch instruction the RegDst signal is “don’t care” because it does not write to any registers. Therefore, the programmer cannot always set RegRt to 0 to test for this fault. As a result, we cannot reliably test for this fault.

